

COMP09024 Unix System Administration

Lecture 4: Processes, Jobs and Scheduling

Duncan Thomson/Hector Marco

UWS

Trimester 1 2020/21

Outline

4.1 Processes

- What is a Process?
- Process Lifecycle
- Examining Processes: `ps`
- The `top` Command
- Examining System Information

4.2 Controlling Processes

- Stopping Interactive Processes
- Signals
- The `kill` Command
- Process Priority and `nice`

4.3 Job Control from the Shell

- Foreground and Background Tasks

- `jobs`, `fg` and `bg`
- Process I/O and Terminals

4.4 Scheduling Regular Tasks

- The `cron` System
- System Tasks and `/etc/crontab`
- User Crontabs
- `cron` and Debian

4.5 One-off Tasks and System Time

- The `at` System
- System Time: `date`
- Keeping Time

4.1 Processes

What is a Process?

- A *process* is an instance of a running program
- If a program is run multiple times (by same or different users) each 'instance' requires a separate process
- In Unix, every process is uniquely identified by a process ID (PID)
- When a program runs, the OS keeps track of various data:
 - The program binary or *image*
 - Memory used (.text, .data, Stack, Heap, etc.)
 - State information (registers, running threads)
 - Resources (files, network connections)
 - Security information (user/group, permissions)
- These data together form a `process`

Process Lifecycle

- A process is *spawned* by another process, known as its parent (except for PID 1 (`init`), which is started by the kernel)
 - The parent process uses the `fork` system call, which makes a copy of itself
 - One of these copies uses the `exec` system call, which replaces itself with the new program
- An active process is usually in one of three states: running, sleeping (waiting for something), or stopped
- A process informs its parent when it terminates
- If its parent isn't listening, it remains a 'zombie'
- If its parent has terminated, it becomes an 'orphan' and is adopted by PID 1 (which becomes its new parent process)

Examining Processes: `ps`

- The `ps` command shows a list of processes
- By default, it shows only:
 - That user's *interactive* processes
 - PID, controlling terminal, time running, and command
- The `ps` command has a wide range of options, which vary depending on the Unix variant
- The GNU version of `ps` common in Linux distributions includes
 - Unix System V options, preceded by `-`
 - BSD options, without `-`
 - GNU options, preceded by `--`

Some Options to `ps`

- Some options changing the processes displayed:
 - `a` show processes from other users
 - `x` show processes without a terminal
 - `-e` (or `ax`): show all processes
 - `r` show only running processes
- Some options changing what information is displayed:
 - `-f` full format (parent PID)
 - `-F` extra full format (PPID and memory usage)
 - `-l` long format (UID, status and priority)
 - `u` user-oriented format (username, CPU and memory usage, status)
- Various flags to `ps` allow much finer control: selection criteria; order and format of fields; sort order; and more — `man ps` gives full details

The `top` Command

- Sometimes it is useful to see a realtime view of the system and processes, rather than the snapshot given by `ps`
- The `top` command shows a continually refreshed view of some key system statistics, and some processes, sorted (by default) by CPU usage
- Statistics include load averages, number of users, process counts, CPU usage statistics, and memory/swap usage
- `top` is exited by pressing `q` (quit)
- A variety of keystrokes allow control of fields displayed, sort order, refresh rate (pressing `?` gives a list)
- Processes may also be reprioritised or killed

Examining System Information

- System information commands vary a lot from one Unix system to another
- The following all work on Linux, some may also work on other Unices
- `uname` gives system name and kernel versions
- `free` gives memory usage; `vmstat` gives more detail
- `df` summarises disk space usage
- `uptime` shows system uptime
- `lscpu` shows basic CPU information (more in `/proc/cpuinfo`)
- `lspci` and `lsusb` show hardware connected to PCI and USB busses

4.2 Controlling Processes

Stopping Interactive Processes

- Interactive processes — those which are (in theory) interacting with a user via a terminal — may be terminated
- Termination of an interactive process is done by pressing **Ctrl-C** on the controlling terminal
- This sends an ‘interrupt’ (`SIGINT`) signal to the process, which normally results in it terminating
- Some other control characters are sometimes confused with:
 - Ctrl-D (End Of File) is often used to terminate input to a process (for example to logout of a system)
 - Ctrl-Z sends a `SIGTSTP` signal, which stops (or suspends) a process; this does not terminate it, but detaches it from the terminal while it waits for a signal to continue

Signals

- Signals facilitate sending a basic notification to a process
- Each type of signal has a number and an associated name
- Other than the signal number, no other data is sent to a process when a signal is sent
- Some important signals for the Unix administrator include:
 - `SIGINT` (2) interrupts the process, usually terminating it
 - `SIGTERM` (15) asks the process to terminate
 - `SIGKILL` (9) kills the process
 - `SIGHUP` (1) for 'hang up' — server processes (daemons) will often reload configuration data on receipt of this signal
 - `SIGTSTP` or `SIGSTOP` stops execution
 - `SIGCONT` restarts execution

The `kill` Command

- The `kill` command sends a signal to a process
- The syntax is `kill -SIGNUM PID`
 - *SIGNUM* is the signal number
 - *PID* is the process ID
 - If the signal number is omitted, the default is `SIGTERM` (15)
- Normal users may only send signals to their own processes
- `root` can send signals to any process
- An alternative to providing a PID is providing a job number (see later) with a `%` sign in front eg: `kill -9 %1`
- Note that `top` allows sending signals using the `k` key

Process Priority, `nice` and `renice`

- The exact priority processes are given depends on scheduling algorithms used by the specific Unix kernel
- The *nice* value of a process is a user interface to influence priority: the nicer, the lower priority
- The default nice value is 0, and values range from -20 to 19
- Processes may be started with positive (or negative, by `root`) nice values using: `nice NICE command args`
- The niceness of a running process can be changed with `renice NICE PID`
- Normal users can only `renice` their own processes, and only increase the nice value; `root` can `renice` arbitrarily
- In `top`, processes can be `reniced` with the `r` key

4.3 Job Control from the Shell

Foreground and Background Tasks

- The shell normally starts processes in the *foreground*
- This means the shell must wait for the process to terminate before anything else can be done
- For non-interactive batch tasks, there is generally no need to wait for such processes to terminate — these can be run in the *background*
- Interactive processes can be run in the background, but
 - Output from the process will be mixed in with output from the foreground task
 - If the process requires input, it will stop and wait
- A process is run in the background by appending `&` to the command line, eg: `sleep 50 &`
- The shell prints a job number (in brackets) and the PID

jobs, fg and bg

- When a number of tasks are running, it can become difficult to keep track of them all
- The `jobs` shell command gives a list of jobs, each with a job number and a status (running or stopped) and an indication of which is the 'current' job (with a + sign)
- Jobs can be moved to the foreground or background with the `fg` and `bg`
- By default these act on the current job, but a job number can be given by preceding it with `%`
- If a job's status changes to stopped (waiting for input) or completed, the shell will report this at the next opportunity (usually between commands)

Process I/O and Terminals

- All Unix processes have standard input, output and error streams,
- For batch processes, these would typically be files, or possibly other processes (we will look at this next week)
- For interactive processes, these are terminal devices
- Originally, terminal devices files represented serial ports on the Unix machine (connected to ‘dumb terminals’)
- Nowadays, terminals usually represent CLI sessions with the machine via non-physical connections (shell windows or network connections)
- Linux supports virtual terminals (via Alt-F1–Alt-F8)
- The `screen` utility is a useful virtual terminal with multiplexing and attach/detach capability from terminals

4.4 Scheduling Regular Tasks

The cron System

- Often batch tasks need to be scheduled on a regular basis
- The `cron` system allows such tasks to be scheduled
- Allows both system tasks and user tasks to be scheduled
- Consists of:
 - A daemon (server process): `crond`
 - A system configuration file: `/etc/crontab`
 - Individual user crontabs in `/var/spool/cron/crontabs`
- Crontabs specify when and what jobs are to be run

System Tasks and `/etc/crontab`

- The `/etc/crontab` file is the system crontab
- Comment lines begin with `#`
- Each line has seven space or TAB-separated fields:
 - The first five are time fields, which are (in order): minute (0–59), hour (0–23), day of month (1–31), month (1–12) and day of week (0–7, or names)
 - Time fields may include the `*` wildcard, lists, or ranges (eg `1–7`, `15–21`)
 - The sixth field is the username to run the job as
 - The seventh field is the command to run

```
# run as root on Sundays at 03:30
30 3 * * 0 root /root/backup.sh
# run on weekdays in April, every 8 hours
0 1,9,17 * 4 1-5 root /root/taxbackup.sh
```

User Crontabs

- Depending on the system setup, users may also be able to schedule jobs with `cron`
- The `/etc/cron.allow` and `/etc/cron.deny` files can specify which users can/cannot schedule jobs
- Users use the `crontab` command with appropriate flag:
 - `-l` lists personal crontab
 - `-e` edits personal crontab
 - `-r` removes personal crontab
- Fields are as for system crontab, but without 'user' field

cron and Debian

- Debian extends the system crontabs in two ways for easy management of regular tasks
- Four directories in `/etc`: `cron.hourly/`, `cron.daily/`, `cron.weekly/` and `cron.monthly/` hold shell scripts which require to be run at the given intervals
- Additional (system) crontabs are included from files in `/etc/cron.d/` — this allows packages to add their own crontab entries easily
- Debian also uses `anacron` by default — a version of `cron` which can be used in environments where the system is not always switched on

4.5 One-off Tasks and System Time

The at System

- The `at` system (under the control of `atd` - the `at` daemon) allows a job or sequence of commands to be run at a specific time
- Jobs are submitted using `at` with a time:
 - `at HH:MM`
 - `at DAY HH:MM`
 - `at YYYY-MM-DD HH:MM`
 - More complex time specifications are possible, eg: `teatime + 3 days`
- Commands can be typed into standard input (ending with Ctrl-D) or supplied in a separate file with the `-f` flag, eg:

```
root@debian:~$ at -f jobfile.sh 20:15
```

More `at` Commands

- The `atq` command shows jobs waiting in the `at` queue
- Each job has a number (in first column of list)
- Jobs can be removed using `atrm` with the job number
- The `batch` command runs the given job when system load is less than 1.5 (set when the `at` system is started)
- `at` has several queues (the `-q` flag can select which one):
 - `a` is the default
 - `b` is for batch jobs
 - `c-z` are used for running with increasing `nice` values
- `/etc/at.allow` and `/etc/at.deny` allow control over which users may use the `at` system

System Time: `date`

- How do `cron` and `at` know what time it is?
- System time can be examined using the `date` command (format may be specified with `+` along with field specifiers):

```
$ date
Thu Jan 2 03:06:21 GMT 2014
$ date +"%Y-%m-%d %H:%M:%S"
2014-01-02 03:08:22
```

- `root` may also set the system time manually with `date`

```
# date 01020310.30
Thu Jan 2 03:10:30 GMT 2014
```

- The `cal` command can be used to display a basic calendar

Keeping Time

- Usually system time is read from hardware clock on boot
- The hardware clock can be accessed with `hwclock`
- The hardware clock is usually running on UTC
- System time is calculated from this using the time zone
- Summer time adjustments are made automatically
- Even better than using the hardware clock is using Network Time Protocol (NTP)
- Normally there are NTP servers within an organisation
- NTP clients can set their time from these
- NTP servers use external servers to remain accurate

Summary

- Processes overview
- `ps` and `top`
- Signals and `kill`
- Process priority, `nice` and `renice`
- Job control: `jobs`, `fg` and `bg`
- Backgrounding tasks with `&`
- Terminating and pausing processes with `Ctrl-C` and `Ctrl-Z`
- Interactive and batch processes and terminal devices
- Scheduling regular tasks with `cron`
- System and user crontabs
- The `at` system
- System time and `date`