

COMP09024 Unix System Administration

Lecture 5: Redirection, Pipes and Filters

Duncan Thomson/Hector Marco

UWS

Trimester 1 2020/21

Outline

5.1 Redirecting I/O

- Process Input and Output
- Standard Streams
- Redirection of Streams

5.2 Stream Redirection

- Redirecting `stdout`
- Redirecting `stdin`
- Redirecting `stderr`
- Pipelines and the Pipe Operator
- Advanced Redirection

5.3 Commands and Exit Status

- Command Strings

- Exit Status
- Conditional Command Strings

5.4 Filters

- Regular Expressions and `grep`
- Selecting with `cut`
- Various Text Filters
- The Stream Editor: `sed`

5.5 More Unix Commands

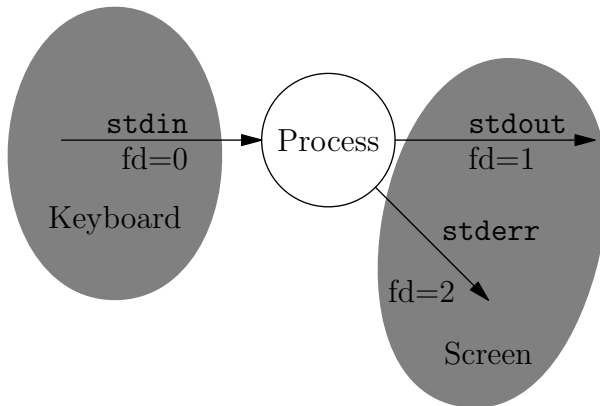
- Finding Files
- Archiving
- Comparing Files
- Calculations

5.1 Redirecting I/O

Process Input and Output

- Most programs expect to:
 - Take some kind of input (maybe from a file)
 - Process the data
 - Send the results to some kind of output (perhaps a file)
- Some processes may also produce error messages
- Files — or other inputs or outputs in Unix (everything is a file!) — are identified by a *file descriptor*
- By default, each processes is started with three *standard streams*:
 - *Standard input* (`stdin`), identified by file descriptor 0
 - *Standard output* (`stdout`), identified by file descriptor 1
 - *Standard error* (`stderr`), identified by file descriptor 2
- By default, these are the keyboard and the screen

Standard Streams



Redirection of Streams


- A key concept of Unix is that these standard streams can be *redirected*
- These can be redirected to/from anything which can be described by a file descriptor, including:
 - Sending `stdout` output to a file
 - Sending `stderr` output to a file
 - Taking `stdin` input from a file
 - Taking `stdin` input from `stdout` (or `stderr`) of another process
 - Sending `stdout` (or `stderr`) output to `stdin` of another process
- Many programs are written specifically to allow chaining together in this way, and are often known as *filters*

5.2 Stream Redirection

Redirecting `stdout` to a File

- To redirect `stdout` to a file, the `>` operator is used
- Example: redirecting the output of `ls` to the file `files.txt`


```
user@debian:~$ ls -l >files.txt
```



Redirecting `stdin` from a File

- To redirect `stdin` from a file, the `<` operator is used
- Example: redirecting the input of `more` from the file `/etc/passwd` (bit of a silly example, since we could just do `more /etc/passwd`)


```
user@debian:~$ more </etc/passwd
```



Redirecting `stderr` to a File

- To redirect `stderr` to a file, the `2>` operator is used
 - Here, the `2` references `stderr`
- Example: redirecting errors from the `cat /etc/*` command to the file `errs.txt`


```
user@debian:~$ cat /etc/* 2>errs.txt
```



The Pipe Operator

- The `|` operator connects `stdout` of the first process to `stdin` of the second
- Example: search for the process name `init` in output from `ps`
 - The `grep` command searches for a string

```
user@debian:~$ ps -ef | grep init
```



Advanced Redirection

- Additional redirections may be possible, eg `bash` allows:
 - Appending `stdout` to end of an existing file with `>>`
 - Redirecting `stderr` into `stdout` with `2>&1`
 - Redirecting `stderr` to another command with `|&`
 - Redirecting both `stdout` and `stderr` to a file with `&>`
 - A 'here-document': taking `stdin` of a command from `stdin` (ended by a specified string, eg 'EOF') with `<<EOF`
- There are also some special device files, for example:
 - `/dev/null` discards anything sent to it
 - `/dev/zero` provides an endless string of zeros
 - `/dev/random` provides (pseudo)-random data
- Possible to combine redirections in a command string, eg:

```
user@debian:~$ ps -ef | grep init | more
user@debian:~$ who | grep alice >alice.txt
```

5.3 Commands and Exit Status

Command Strings

- Sequential execution of commands one after another (`command2` is run after `command1` has completed)

```
command1 ; command2
```

- Parallel execution of commands (`command1` in background):

```
command1 & command2
```

- Pipelined commands (moving `stdout` of `command1` to `stdin` of `command2`):

```
command 1 | command2
```

Exit Status

- Every process when it terminate has an *exit status* — an integer indicating the outcome of the process
- This can be accessed using the variable `$?`
- The value is usually 0 for ‘success’, and may take other values for other occurrences
- For example:

```
user@debian:~$ grep -sq root /etc/passwd ; echo $?  
0  
user@debian:~$ grep -sq raat /etc/passwd ; echo $?  
1  
user@debian:~$ grep -sq root /etc/pass ; echo $?  
2
```

Conditional Command Strings

- There are also two ways of combining commands 'conditionally' — depending on the exit status of the first
- 'AND' command combination uses `&&`:
 - `command1 && command2`
 - `command2` is only run if `command1` has a zero exit status
- 'OR' command combination uses `||`:
 - `command1 || command2`
 - `command2` is only run if `command1` has a non-zero exit status
- Overall exit status in both cases is status of last command run

5.4 Filters

Filters

- A large number of Unix tools are designed to operate by taking data on standard input, transforming it in some way, and sending it out to standard output
- Such programs are known as *filters*
- Generally follow the Unix philosophy of ‘do one thing well’
- We have already come across a couple:
 - `cat` — takes standard input (or a list of files) and concatenates them onto standard output
 - `grep` — searches for a string in standard input, prints matches on standard output (we’ll examine this more closely in a moment)
- We’ll now look at some Unix filters

Searching with `grep`

- `grep` (stands for General Regular Expression Parser) searches for strings (or *regular expressions*) in standard input (or files)
- Matching lines are sent to standard output
- A wide range of options are available, including:
 - `-v` prints non-matching lines
 - `-i` matches case-insensitively
 - `-F` or `-E` match fixed strings or *extended* regexps
 - `-l` (or `-L`) only list filenames with (or without) matches
 - `-c` only output count of matching lines (per file)
 - `-q` (and `-s`) suppress output (or error messages)
 - `-A` and `-B` include context after and before matches

Regular Expressions

- (Basic or extended) *regular expressions* (regexp or regex) are used for matching patterns in strings using special purpose characters (see `man grep`)
- `.` matches any single character
- `[chars]` matches any character listed (which may include ranges like `[a-z]` or `[0-9]`)
- `[^chars]` matches any character except those listed
- `^` and `$` match beginning and end of line respectively
- `?` matches 0 or 1 of preceding character
- `*` matches 0 or more of preceding character
- `+` matches 1 or more of preceding character
- `{n}` matches exactly *n* of preceding character
- `(re1|re2)` matches either of the subexpressions

Examples of Regular Expressions

- `abcde` — matches the string `abcde`
- `abcd.` — matches any 5-char string starting with `abcd`
- `Thomp?son` — matches either `Thomson` or `Thompson`
- `Gra(eme|ham)` — matches either `Graeme` or `Graham`
- `[0-9A-F]` — matches any (uppercase) hexadecimal digit
- `^[^a-z]` matches any line not starting with a lower case letter
- `[Ww]ee+` — matches `'Weeee'`, `'wee'`, `'weeeeeeeee'` (but not `'we'` `'woooo'` or `'Wii'`)
- `07[0-9]{9}` — matches a UK mobile phone number
- `[A-Z][A-Z]?[1-9][0-9]?[0-9][A-Z]{2}` — matches UK postcode

Examples using `grep`

Some examples using `grep`:

```
user@debian:~$ cat file.txt
s1 newa e1
s2 newb e2
s3 newc e3
s4 newd e4
user@debian:~$ grep new[ac] file.txt
s1 newa e1
s3 newc e3
user@debian:~$ grep new[a-c] file.txt
s1 newa e1
s2 newb e2
s3 newc e3
user@debian:~$ grep -E "(e1|e2)" file.txt
s1 newa e1
s2 newb e2
```

Selecting text with `cut`

- The `cut` command picks specific characters or fields
- It either operate in byte, character or field mode:
- In character mode (`-c`), a list of characters is provided:
 - Ranges are specified by using the `-` symbol
 - More than one range may be given, spearated by commas
 - Example: `dpkg -l | cut -c5-40` shows a list of package names (only) installed on Debian
- In field mode, (`-f`) a list of fields is given
 - Field separator is TAB or provided with `-d`
 - Several ranges may be provided, as with `-c`
 - Example: `cut -d: -f1,5 /etc/passwd` shows usernames and with full user's name and contact details

sort and uniq

- `sort` sorts lines in `stdin` (or a file)
 - `-a` performs an alphabetical sort
 - `-g` performs a general numerical sort
 - `-R` sorts randomly
 - `-M` sorts by month name
 - `-r` reverses order of sort
 - `-k` specifies which 'key' (field) to sort on
 - `-t` specifies field separator character
 - Example: `sort -t: -k3 -g /etc/passwd` — sorts `/etc/passwd` numerically by UID field
- `uniq` searches for unique or duplicate lines in sorted input (and removes duplicates by default)
 - `-d` only prints duplicates
 - `-u` only prints unique lines

head, tail and tac

- The `head` command by default prints the first 10 lines of `stdin` or a file
 - The number of lines can be given with `-`
- The `tail` command by default prints the last 10 lines of `stdin` or a file
 - The number of lines can be given with `-`
 - The `-f` allows ‘following’ of a file as it is added to
- The `tac` command prints the lines of `stdin` (or file) in reverse order
 - `-s` specifies a separator other than newlines
- Example: `head /etc/passwd | tail -5` prints lines 6-10 of the `/etc/passwd` file

`wc` and `tee` and `tr`

- The `wc` command counts characters, words and/or lines
 - `-c` prints a count of characters
 - `-w` prints a count of words
 - `-l` prints a count of lines
- The `tee` command copies `stdin` to `stdout`, sending a copy into a file specified (like a T-piece in a pipe)
- `tr` allows translating from one set of character to another
 - `-d` deletes specified characters instead
- **Example:** `tr a-z A-Z <file.txt | tee upper.txt | wc -w` puts an uppercase version of `file.txt` into `upper.txt` and prints a word count of the result

- `sed` allows arbitrary edits to be made to data
- A very powerful tool, but not easy to understand
- Commands can be provided to it either
 - Using `-e` to give a command on the command line
 - Using `-f` to specify a file containing commands
- Each command consists of:
 - An optional address or address range specifying which lines to edit
 - A command, which is executed for all matching lines
- A `sed` script may include multiple commands
- An example `sed` command: `sed -e "/hello/y/a-z/A-Z/"` — this capitalises all letters on lines which contain the word 'hello'

sed Addresses and Commands

- Some methods of specifying lines include:
 - Line number, eg `10`
 - Line number range with a comma eg `10,20`
 - A regular expression using `/` as delimiters, eg `/^[0-9]/` will match all lines starting with a digit
- Commands which can be executed include:
 - `a` appends text after matching lines
 - `i` inserts text before matching lines
 - `c` changes lines for new text
 - `d` delete lines
 - `s/regexp/string/` substitutes *string* in place of *regexp*
 - `y/chars/chars2/` substitutes characters in *chars2* for those in *chars1* (like `tr`)

5.5 More Unix Commands

Finding Files with `find`

- The `find` command can find files matching specific criteria
- The syntax is: `find path criteria action`
- The path is the starting point in the filesystem
- Criteria can be combined, and include:
 - `-name` to match on a filename
 - `-atime`, `-ctime` or `-mtime` to match access, change or modification times ('equal', 'after' or 'before' a number of days; there are other variations)
 - `-user` or `-uid` to match the owner
 - `-group` or `-gid` to match a group owner
 - `-perm` to match on file permissions
 - `-size` to match on file size
 - `-type` to match by type (directory, link, etc)
- Actions can include `-print` or executing commands with `-exec`

Archiving with `tar`

- The `tar` command allows creation of archives of multiple files (from Tape ARchive)
- `f` options specifies archive file (almost always used)
- Other options include:
 - `c` to create an archive (from the list of files given)
 - `x` to extract files from an archive
 - `t` to print a table of contents
- `tar` archive files normally end in `.tar`
- `tar` is only one of many archival programs, which include `dd`, `cpio` and others depending on system
- Example: `tar xzf archive.tgz` uncompresses and extracts all files from the file `archive.tgz`

Compressing Files: `gzip`

- A number of compression utilities are available
- `gzip` is one of the most popular
- Default is to compress the given file (giving it a `.gz` extension)
- Level of compression can be controlled with `-1` to `-9`
- Compressed files are uncompressed with `gunzip`
- There are also some special variants of commands which operate directly on compressed files, eg `zcat`, `zmore`, `zgrep`
- Some systems may have other compression commands, for example:
 - `compress` and `uncompress` (the original Unix one)
 - `bzip` and `bunzip`

Comparing Files

- A number of commands are available to compare files
- `diff` is the most widely used — reports on differences between text files
- Can provide context lines
- Is widely used to produce ‘patch’ files (which can be used to update source code using the `patch` command)
- The `cmp` utility also compares files

Calculations

- There are two command line calculators widely available in Unix
- `dc` is a reverse Polish notation (stack-based) calculator
- `bc` is more convenient-to-use calculator
- Both of these provide the user with command-line interface to the calculator
- `dc` also allows providing a script directly on the command line
- The `expr` command is more useful for shorter calculations such as might be used in shell scripts (next week)
- Example: `expr 2 + 2` gives the answer 4

Summary

- Process input and output
- Standard process streams: `stdin`, `stdout` and `stderr`
- Redirection
- File redirection: `>`, `<`, `>>`, `2>` and `<<`
- Pipelines and the pipe operator `|`
- Command exit status
- Chaining commands: `;`, `&`, `&&` and `||`
- `grep` and regular expressions
- Filters: `cut`, `head`, `tail`, `tee`, `sort`, `sed`, `tr` and others
- Finding files matching criteria with `find`
- Archiving and compressing: `tar` and `gzip`
- Finding differences in files: `diff` and `patch`
- Calculations with `expr`, `dc` and `bc`