

# COMP09024 Unix System Administration

## Lecture 8: System Configuration and Monitoring

Duncan Thomson/Hector Marco

UWS

Trimester 1 2020/21

# Outline

## 8.1 Managing Unix Systems

- Command Line
- Configuration Files
- Services
- Secure Shell

## 8.2 The vi Text Editor

- vi Revisited
- Movement Commands
- Editing Commands
- Other Commands
- Example Commands

## 8.3 The Boot Process

- Overview
- Boot Loaders
- Service managers

- System V init and /etc/inittab

- System V init Scripts

## 8.4 Controlling Services and Systems

- Using init Scripts and service
- Special Cases
- Kernel Configuration

## 8.5 Logging, Monitoring and Auditing

- Logging and Log Files
- syslog
- Monitoring Systems
- Auditing Systems

## 8.1 Managing Unix Systems

# Command Line Interface

- Unix servers were originally command line only, and so configuration was performed on the command line
- Although most Unix systems now include a GUI (usually based on X11), command-line text-based administration has persisted
- Advantages include:
  - Existing tools used to perform configuration
  - Low overhead and impact on other software
  - Remote administration identical and low bandwidth
- Disadvantages:
  - System administrators need to be comfortable with CLI and working with text files
- ...but hopefully by now, you already are!?

# Configuration File

- Configuration files are found in `/etc/` (eg: `/etc/passwd`, `/etc/hostname`, `/etc/inittab` ...)
- In many cases, subdirectories of `/etc/` are used for more complex subsystems (eg `/etc/apt/`, `/etc/cups/`, `/etc/network/`)
- In (almost) all cases, configuration files are text-based
- These may be documented in three ways:
  - A `man` page (in section 5)
  - Comments in the file itself
  - Example configurations (sometimes included in the comments)
- Since they are text-based, text editors (like `vi`) can be used to make changes

# Servers and Daemons

- Many services on Unix run without being connected to a terminal
- Such server processes are often known as *daemons*
- Such processes need to:
  - Be automatically run at boot time (if required)
  - Be able to be stopped and started as required
  - Be able to be told to re-read their configuration files if they are changed
  - Be monitored and restarted if they stop unexpectedly
- The `init` system typically provides mechanisms for some or all of these

# Secure Shell

- If systems are administered remotely, a mechanism to access the command line remotely is required
- In days gone by, `telnet` or the `rlogin`:
  - Authentication but not Encryption (data unencrypted)
- Modern Unix systems provide Secure Shell (SSH):
  - Authentication and Encryption
- A secure shell server is run on the server (`sshd`).
- The `ssh` command will connect to a remote machine running SSH server, and once credentials have been checked command line access is granted

```
ssh servername
```

```
ssh username@servername
```

# Secure Shell Capabilities

- SSH client checks server key on connection, and refuses to connect if this has changed (to prevent man-in-the-middle attack).
- SSH can copy files to/from the remote server (with `scp`):  

```
scp localfile user@server:remotefile  
scp user@server:dir/remotefile localfile
```
- SSH provides key management to automate authentication
- SSH can tunnel X11 (GUI) data with the `-X` flag
- SSH can tunnel TCP connections through the connection:
  - `-R` to tunnel from remote port to local port
  - `-L` to tunnel from local port to remote port



## 8.2 The `vi` Text Editor

# `vi` Revisited

- `vi` is the most ubiquitous text editor on Unix
- Also the basis for `sed` the stream editor
- ... and is also quick to use, once you've learned the basics
- Remember, `vi` is modeful:
  - In command mode (or 'normal' mode), every keystroke is a (case-sensitive) command
  - In insert mode, you are typing into the file (exit with ESC)
  - In 'bottom line mode', commands from the legacy `ex` editor can be used (including file save, quit, etc)
- It's also *case-sensitive* (upper and lower case commands are often related, but not the same)
- All commands can be preceded by a number, which in most cases repeats the command that number of times

# Movement Commands

- `h j k l` — movement left, down, up, right (or use arrows)
- `w, b, e, W, B, E` — move forwards, backwards or to end of a word (including punctuation with capitals)
- `0` and `$` — move to start or end of line
- `H, M` and `L` — move to top, middle or bottom of screen
- `Ctrl-B` and `Ctrl-F` — move back / forward a screenful
- `Ctrl-U` and `Ctrl-D` — move back / forward half-screenful
- `f c` or `F c` — move onto next/previous occurrence of char
- `t c` or `T c` — move up to (but not including) next/previous occurrence of character
- `nG` — go to a given line number (or to end of file without `n`)
- `%` — go to matching opening or closing parenthesis
- `/ regex` or `? regex` — search forward / backward

## Editing Commands

- `a` or `A` — append after cursor / end of line (insert mode)
- `i` or `I` — insert before cursor / start of line (insert mode)
- `o` or `O` — open new line after/before current one (insert)
- `dmove` or `ymove` — delete or copy (yank) specified text
- `cmove` — change specified text (insert mode)
- `D` or `C` — delete or change to end of line (insert mode)
- `dd`, `yy` or `cc` — delete, yank or change whole line
- `p` or `P` paste deleted/yanked content after / before current position
- `x` or `X` — delete character forwards / backwards
- `rc` — replace current character
- `~` — change case of current character
- `J` — joins current and next line

## Other Commands

- `u` — undo last command
- `.` — repeat last editing command
- `n` or `N` — repeat last regexp search (same / reverse dir)
- `:w` or `w filename` — save (write) file
- `:q` or `q!` quit / quit without saving
- `:wq` — save and quit
- `:r filename` — read in a file at cursor position
- `:s/regexp/string/` — search and replace first occurrence on current line, or precede with a line range:
  - `10,20` for lines 10–20
  - `1,$` for all lines
  - `.,+10` for current line to 10 lines later
  - Suffixed `g` replaces all occurrences on the line(s)

# Example Commands

- The true power of `vi` comes from:
  - Combining edit and movement commands
  - Preceding commands with numbers
  - Using the ‘repeating’ commands
- Examples:
  - `dG` — delete from here to the end of the file
  - `10dd` — delete 10 lines (which can be pasted back with `p` elsewhere)
  - `20~` — change case of 20 characters
  - `c3w` — change the following 3 words into inserted text
  - `d%` — delete contents of this parenthesis

## 8.3 The Boot Process

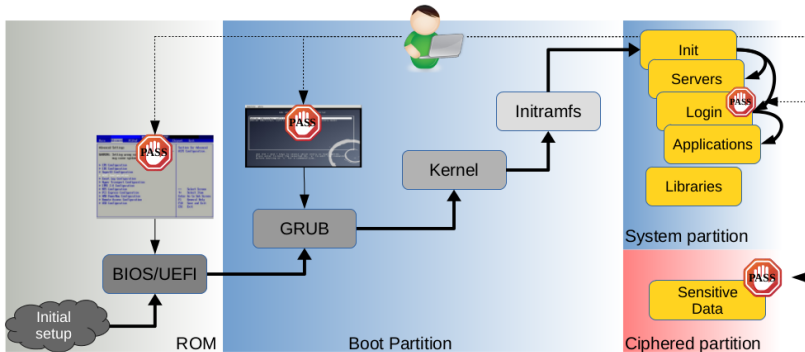
# Overview of System Boot (1/2)

When a system boots it goes through a number of stages:

- ➊ Load and run initial firmware (eg BIOS, UEFI, etc)
  - Perform power-on self-test (POST)
  - Detect attached devices
  - Locate boot loader (eg on master boot record (MBR))
- ➋ Load and run boot loader software (eg GRUB, LILO, etc)
  - Read devices and filesystems
  - Locate kernel and any other initial filesystem
- ➌ Load and run OS kernel
  - In some cases an initial filesystem may also be required
- ➍ Mount main root filesystem at `/`
- ➎ Run first process (`init`)
- ➏ Bring up other filesystems and processes as required (under control of `init`)



# Overview of System Boot (2/2)



Boot system Overview. User's HOME partition ciphred.

(source: <http://hmarco.org/bugs/CVE-2015-8370-Grub2-authentication-bypass.html>)

# Boot Loaders

- A range of boot loaders are available for Unix systems
- Some are tied strongly to one operating system, e.g:
  - Linux LOader (LILO)
  - `isl` for HP-UX
- GNU Grand Unified Bootloader (GRUB) is a modern modular boot loader used by multiple Unix variants, including Solaris and Linux
  - Configured from script in `/boot/grub/grub.cfg`
  - Can load modules for filesystems, etc
  - Allows selection between multiple boot options, or manual boot

# Service managers

## System-V, Upstart and systemd:

- The classical “System-V” `init` service.
  - An obsolete way to start and stop groups of services.
  - It uses the runlevel concept.
- Around 2006, the “Upstart” `init` service replaced SysV.
  - It replaces “System-V”.
  - No more `/etc/inittab` file.
- Nowadays, most Linux distributions use “systemd”:
  - The runlevel concept is obsolete.
  - It provides a mapping between runlevels and `systemd` targets (e.g: runlevel 5 → `graphical.target`)

# System V `init`

- The traditional Unix `init` system is System V `init`
- This defines a number of runlevels:
  - Runlevel 0 is for shutting down the system
  - Runlevel s is for single-user mode (for maintenance)
  - Runlevel 6 is normally for rebooting
  - Runlevels 1–5 can be manually configured
- Overall behaviour of `init` controlled by `/etc/inittab`
- Scripts used to start/stop services when changing runlevel
- The `telinit` command can switch between runlevels
- Other commands provide shortcuts for certain runlevels:
  - `reboot` reboots the system immediately (runlevel 6)
  - `halt` shuts down the system immediately (runlevel 0)
  - `shutdown` allows delayed halt (or reboot with `-r`) along with a message for users

# /etc/inittab

- `/etc/inittab` has four colon-separated fields
- ID field identifies entry with up to four characters
  - Some IDs are traditionally used for certain entries, eg `id` for initial default runlevel; numbers relating to login processes on `ttys`; etc
- Runlevels field lists runlevels in which service should be active
- Action specifies how to run process, eg:
  - `initdefault` specifies default runlevel
  - `sysinit` and `boot` run on boot only
  - `once` and `wait` run once on entering runlevel
  - `respawn` restarts on exit
  - `ctrlaltdel` runs when Ctrl-Alt-Del pressed
- Process specifies command (and arguments) to run

## Example `/etc/inittab`

```
# default runlevel is 2
id:2:initdefault:

# system initialisation on boot
si::sysinit:/etc/init.d/rcS

# single-user mode
~~:S:wait:/sbin/sulogin

# rc script for each runlevel
10:0:wait:/etc/init.d/rc 0
11:1:wait:/etc/init.d/rc 1
12:2:wait:/etc/init.d/rc 2
13:3:wait:/etc/init.d/rc 3
14:4:wait:/etc/init.d/rc 4
15:5:wait:/etc/init.d/rc 5
16:6:wait:/etc/init.d/rc 6
```

```
# what to do what Ctrl-Alt-Del
ca:12345:ctrlaltdel:/sbin/shutdown \
    -t1 -a -r now

# what to do when power fails
pf::powerwait:/etc/init.d/powerfail \
    start
pn::powerfailnow:/etc/init.d/powerfail
    now
po::powerokwait:/etc/init.d/powerfail
    stop

# login prompts on virtual terminals
1:2345:respawn:/sbin/getty 38400 tty1
2:23:respawn:/sbin/getty 38400 tty2
3:23:respawn:/sbin/getty 38400 tty3
4:23:respawn:/sbin/getty 38400 tty4
```

# init Scripts

- The `rc` command (called in the previous example) runs so-called `init` scripts for various services
- The `init` scripts are stored in `/etc/init.d/` or similar, eg `/etc/init.d/cron`
- Each script can take various parameters
  - The `start` parameter starts the service
  - The `stop` parameter stops the service
  - `restart` and/or `reload` parameters may also work
- The `service` command can often be used instead
- For example, to stop the `cron` service, can use either:

```
/etc/init.d/cron stop  
service cron stop
```

# Systemd `init`

- It is a replacement of the UNIX System V.
- Unification of basic Linux configurations and service behaviors across all Linux distributions.
- Instead of `/etc/inittab` we have a `/etc/systemd/system/` directory. This directory holds symlinks to `/lib/systemd/system/` which contains `init` scripts.
- Unlike other `init` systems, you do not have to know a scripting language to interpret the `init` files used to boot services or the system.
- `systemctl [start,stop,restart,reload,status] [name.service]`
  - `systemctl status boot.mount`



# Systemd `init`

## CUPS status example (an open source printing system):

```
$ systemctl status cups.service
```

```
hecargi@macal:system$ systemctl status cups.service
● cups.service - CUPS Scheduler
   Loaded: loaded (/lib/systemd/system/cups.service; enabled; vendor preset: enabled)
   Active: active (running) since Wed 2018-10-17 17:48:54 BST; 57min ago
     Docs: man:cupsd(8)
  Main PID: 1193 (cupsd)
    Tasks: 1 (limit: 4915)
   CGroup: /system.slice/cups.service
           └─1193 /usr/sbin/cupsd -l

Oct 17 17:48:54 macal systemd[1]: Started CUPS Scheduler.
```

## Systemd services:

```
$ ls -l /lib/systemd/system/*.service
```

```
-rw-r--r-- ... /lib/systemd/system/accounts-daemon.service
-rw-r--r-- ... /lib/systemd/system/acpid.service
-rw-r--r-- ... /lib/systemd/system/alsa-restore.service
-rw-r--r-- ... /lib/systemd/system/alsa-state.service
...      ...      ...
```

# Systemd - Linux distributions

Linux distribution ↕	Date added to <b>software repository</b> <sup>[a]</sup> ↕	Enabled by default? ↕
Alpine Linux	N/A (not in repository)	No
Android	N/A (not in repository)	No
Arch Linux	January 2012 <sup>[52]</sup>	Yes
CentOS	April 2014	Yes
CoreOS	July 2013	Yes
Debian	April 2012 <sup>[57]</sup>	Yes
Fedora	November 2010 (v14) <sup>[60]</sup>	Yes
Gentoo Linux <sup>[b]</sup>	July 2011 <sup>[61][63][64]</sup>	No
Knoppix	N/A	No <sup>[65][66]</sup>
Mageia	January 2011 (v1.0) <sup>[67]</sup>	Yes
Mint	June 2016 (v18.0)	Yes
openSUSE	March 2011 (v11.4) <sup>[69]</sup>	Yes
Red Hat Enterprise Linux	June 2014 (v7.0) <sup>[71]</sup>	Yes
Slackware	N/A (not in repository)	No
Solus	N/A	Yes
SUSE Linux Enterprise Server	October 2014 (v12)	Yes
Ubuntu	April 2013 (v13.04)	Yes
Void Linux	June 2011, removed June 2015 <sup>[73]</sup>	No

## `/etc/rcN.d/`

- The particular `init` scripts started (and stopped) in particular runlevels are controlled by the contents of the `/etc/rcN.d/` directories
- These directories (one per runlevel) contain symbolic links
- Each symbolic link links to an `init` script in `/etc/init.d/`
- Each link is named `LNNservice` (`L` is a letter, `NN` is a number), and the name determines how and when each script is run:
  - Links beginning with `K` (for kill) are run with `stop` parameter
  - Links beginning with `S` are run with the `start` parameter
  - Scripts are run in order of the number `NN`

## 8.4 Controlling Services and Systems

# Using `init` Scripts and `service`

- If System V `init` is used, the canonical way to control services is via the `init` script:
  - `/etc/init.d/servicename start` to start a service
  - `/etc/init.d/servicename stop` to stop a service
  - `/etc/init.d/servicename restart` to restart a service
  - `/etc/init.d/servicename reload` to reload configuration for a service
- Usually the `service` command provides another way of accessing this functionality:
  - `service servicename start` to start a service
  - `service servicename stop` to stop a service
  - `service servicename reload` to reload config
  - ...and so on

## Special Cases

- Some services provide their own commands to start, stop and otherwise control operation, examples include:
  - The Apache webserver (`apachectl` or `apache2ctl` commands)
  - The BIND DNS server (the `rndc` command)
- These are usually additional to System V `init` scripts, rather than instead of them
- Other `init` systems provide their own commands:
  - `systemd` **uses** `systemctl`
  - `launchd` **uses** `launchctl`
  - Upstart **uses** `initctl`

# Kernel Configuration I

Most Unix systems have at least two mechanisms for configuring how the kernel operates:

## 1 Add kernel features:

- Load new modules to the running kernel (e.g. drivers)
- In Linux, all modules end with the `.ko` extension
- Modules can be found in `/lib/modules`
- Linux version 5 modules are in  
`/lib/modules/5.0.0-23-generic/kernel/`
- `modprobe ntfs` adds NTFS file system support
- `lsmod` to list currently loaded modules

```
user@debian~$ lsmod
```

Module	Size	Used by
ntfs	106496	0
pci_stub	16384	1
...	...	...

# Kernel Configuration II

## 2 Configure kernel options

- Controlling the run-time behaviour of the kernel
- `sysctl -a` lists all `sysctl` settings
- `sysctl sysctl.name` prints individual setting
- `sysctl sysctl.name = value` sets a value
- `/proc/sys/kernel` contains files controlling a range of kernel parameters

Example: Obtain the current Address Space Layout Randomization (ASLR) mode:

```
user@debian~$ sysctl kernel.randomize_va_space
kernel.randomize_va_space = 2
user@debian~$ cat /proc/sys/kernel/randomize_va_space
2
```

- 0 ASLR is off
- 1 ASLR is on but HEAP is not randomized
- 2 ASLR is on and HEAP is randomized



## 8.5 Logging, Monitoring and Auditing

# Logging and Log Files

- When a system or service is not running correctly, it can be difficult to diagnose the problem using only the `stderr` command-line output
- Most services can produce detailed log messages
- These are typically stored in text-based ‘log files’, usually found under `/var/log/`, eg:
  - `/var/log/syslog` for general messages
  - `/var/log/auth` for authentication logs
  - `/var/log/wtmp` contains user login records (not text)
  - `/var/log/apache2/access.log` for web access logs
- Some services write directly to log files, others use the `syslog` system to unify and control their log messages

# syslog

- The `syslog` system — implemented by (eg) `rsyslogd` or `Syslog-NG` — provides a unified logging framework
- Although origins lie in Unix, it is used on many systems
- Every log message contains various pieces of information:
  - Facility (`kern`, `user`, `mail`, `auth`, `daemon`, `local0`–`local7`, ...)
  - Severity (`emerg`, `alert`, `crit`, `error`, `warn`, `notice`, `info`, `debug`)
  - A timestamp
  - A host name or address
  - A tag (usually the name of program generating the entry)
  - A (human-readable) message
- The `logger` command can send a log message manually
- All messages are sent to `syslog` process, which decides how to deal with each: ignore, append to a file, send over the network, alert one or more users, etc

## Example /etc/rsyslogd.conf File

```
# all security stuff to auth.log
auth.*                /var/log/auth.log
# everything except security to syslog file
*.*;auth.none        /var/log/syslog
# all daemon stuff to daemon.log
daemon.*             /var/log/daemon.log
# mail messages to separate mail log file
mail.warn            /var/log/mail.log
# debug messages only to debug log
*.=debug             /var/log/debug
# critical or emerg messages to everyone logged in
#                    and by UDP to remote server at 10.0.0.10
*.crit               * @10.0.0.10
```

# Monitoring Systems

- One method is to read log files — tedious and inefficient
- `tail -f` can keep track of recent additions to a single file
- When searching for specific information, Unix's filters (such as `grep`, `cut`, `tail` and others) can be very useful,
- There are also tools such as:
  - `ps` and `top` for monitoring processes
  - `who` and `w` for monitoring users
  - `df` and `free` for monitoring filesystem and memory usage
  - `netstat` for network statistics and connections
- Many kernels have additional instrumentation allowing fine-grained information on system performance, eg:
  - `vmstat` on Linux
  - `sysstat` package with `iostat`, `mpstat`, ...

## Other Monitoring Tools

- SNMP: Simple Network Management Tool
  - Allows network access to system performance data (including CPU, memory, disk space, network statistics) from a network management station
  - Requires installing an SNMP *agent* on the managed system
  - Basic tools available for free, or can form part of large network-wide systems (OpenNMS, Tivoli, OpenView, eHealth)
- MRTG (Multi-Router Traffic Grapher)
  - Can provide graphs of performance data over varying time periods
  - Originally designed for network activity, but extends to other system data
- Nagios is another system/network monitoring tool (including SNMP data and other services)

# Auditing Systems

- May sometimes require to audit systems for security — a variety of tools are available
- `nmap` (as seen in *The Matrix*) is a network port scanner with OS fingerprinting and service/software identification
- `tripwire` allows auditing of sensitive system files by creating cryptographic checksums and checking for changes
- Linux includes the 'audit' system which allows detailed monitoring of specified events relating to any system calls, such as:
  - File access
  - Network activity
  - User access

## Summary

- Typical Unix system management: remote CLI, text files
- Secure Shell for remote access and more
- Using the `vi` editor
- Boot process: boot loaders
- The `init` process (PID 1)
- Service managers: System-V, Upstart and `systemd`
- Kernel configuration: modules and `sysctl`s
- Logging and log files
- The `syslog` system: severities, facilities
- Monitoring tools and systems: processes, users, memory; MRTG, SNMP
- Auditing tools: `nmap`, `tripwire`, the Audit system